

Importing the Dependencies

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

Importing the Boston House Price Dataset

```
In [2]: house_price_dataset = sklearn.datasets.load_boston()
```

```
In [3]: print(house_price_dataset)
```

```
{'data': array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e
+02,
    4.9800e+00],
 [2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,
    9.1400e+00],
 [2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,
    4.0300e+00],
 ...,
 [6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
    5.6400e+00],
 [1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,
    6.4800e+00],
 [4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
    7.8800e+00]]), 'target': array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 2
2.9, 27.1, 16.5, 18.9, 15. ,
    18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
    15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
    13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
    21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
    35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
    19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
    20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
    23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
    33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
    21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
    20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
    23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
    15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
    17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
    25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
    23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
    32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
    34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
    20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
    26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
    31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
    22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
    42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
    36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
    32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
    20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
    20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
    22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
    21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
    19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
    32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
    18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
    16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
    13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
    7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
    12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
    27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
    8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
    9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
    10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
    15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
    19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
```

```

29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9]], 'f
eature_names': array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DI
S', 'RAD',
'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7'), 'DESCR': ".. _boston_da
taset:\n\nBoston house prices dataset\n-----\n\n**Data
Set Characteristics:** \n\n :Number of Instances: 506 \n\n :Number of
Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is
usually the target.\n\n :Attribute Information (in order):\n - CRIM
per capita crime rate by town\n - ZN proportion of residential l
and zoned for lots over 25,000 sq.ft.\n - INDUS proportion of non-r
etail business acres per town\n - CHAS Charles River dummy variabl
e (= 1 if tract bounds river; 0 otherwise)\n - NOX nitric oxides
concentration (parts per 10 million)\n - RM average number of ro
oms per dwelling\n - AGE proportion of owner-occupied units built
prior to 1940\n - DIS weighted distances to five Boston employmen
t centres\n - RAD index of accessibility to radial highways\n
- TAX full-value property-tax rate per $10,000\n - PTRATIO pupil
-teacher ratio by town\n - B 1000(Bk - 0.63)^2 where Bk is the
proportion of blacks by town\n - LSTAT % lower status of the popula
tion\n - MEDV Median value of owner-occupied homes in $1000's\n\n
:Missing Attribute Values: None\n\n :Creator: Harrison, D. and Rubinfeld,
D.L.\n\nThis is a copy of UCI ML housing dataset.\nhttps://archive.ics.uci.ed
u/ml/machine-learning-databases/housing/\n\n\nThis dataset was taken from the
StatLib library which is maintained at Carnegie Mellon University.\n\nThe Bos
ton house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic\nprices and
the demand for clean air', J. Environ. Economics & Management,\nvol.5, 81-10
2, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics\n...', Wile
y, 1980. N.B. Various transformations are used in the table on\npages 244-2
61 of the latter.\n\nThe Boston house-price data has been used in many machin
e learning papers that address regression\nproblems. \n \n.. topic:: Re
ferences\n\n - Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying
Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.\n - Qu
inlan,R. (1993). Combining Instance-Based and Model-Based Learning. In Procee
dings on the Tenth International Conference of Machine Learning, 236-243, Uni
versity of Massachusetts, Amherst. Morgan Kaufmann.\n", 'filename': 'C:\\\\Prog
ramData\\Anaconda3\\lib\\site-packages\\sklearn\\datasets\\data\\boston_house
_prices.csv'}

```

```

In [4]: # Loading the dataset to a Pandas DataFrame
house_price_dataframe = pd.DataFrame(house_price_dataset.data, columns = house
_price_dataset.feature_names)

```

```
In [5]: # Print First 5 rows of our DataFrame
house_price_dataframe.head()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST.
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.

```
In [6]: # add the target (price) column to the DataFrame
house_price_dataframe['price'] = house_price_dataset.target
```

```
In [7]: house_price_dataframe.head()
```

Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST.
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.

```
In [8]: # checking the number of rows and Columns in the data frame
house_price_dataframe.shape
```

Out[8]: (506, 14)

```
In [9]: # check for missing values
house_price_dataframe.isnull().sum()
```

```
Out[9]: CRIM      0
ZN          0
INDUS      0
CHAS       0
NOX        0
RM         0
AGE        0
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT      0
price      0
dtype: int64
```

```
In [10]: # statistical measures of the dataset
house_price_dataframe.describe()
```

Out[10]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12

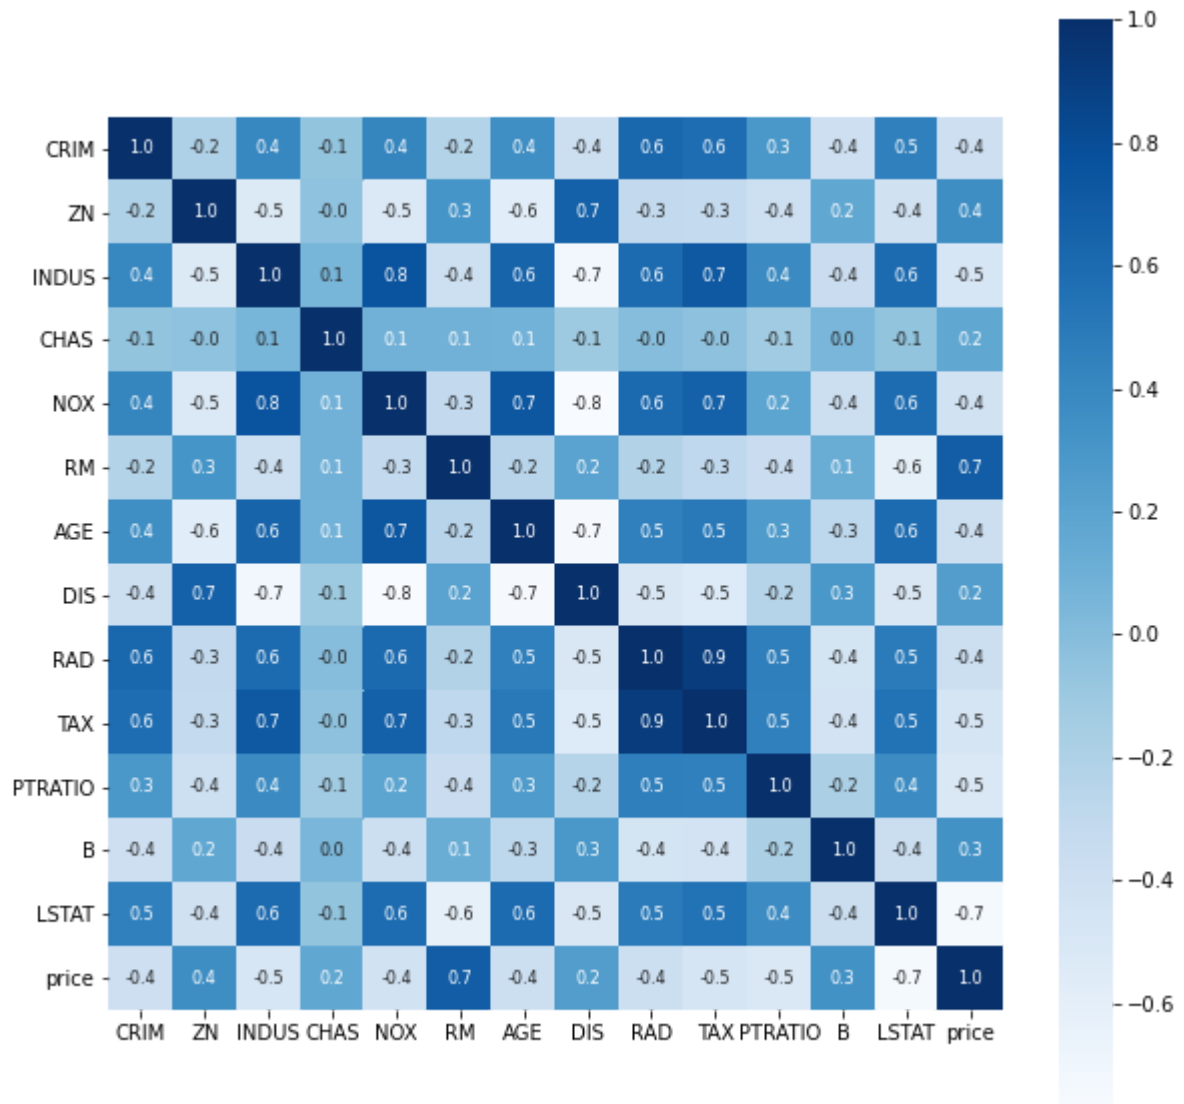
Understanding the correlation between various features in the dataset

1. Positive Correlation
2. Negative Correlation

```
In [11]: correlation = house_price_dataframe.corr()
```

```
In [12]: # constructing a heatmap to understand the correlation
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_
kws={'size':8}, cmap='Blues')
```

Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1c0fe3dc7c0>



Splitting the data and Target

```
In [13]: X = house_price_dataframe.drop(['price'], axis=1)
Y = house_price_dataframe['price']
```

```
In [14]: print(X)
         print(Y)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	
..	
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33
..
501	21.0	391.99	9.67
502	21.0	396.90	9.08
503	21.0	396.90	5.64
504	21.0	393.45	6.48
505	21.0	396.90	7.88

[506 rows x 13 columns]

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2

..	...
501	22.4
502	20.6
503	23.9
504	22.0
505	11.9

Name: price, Length: 506, dtype: float64

Splitting the data into Training data and Test data

```
In [15]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state = 2)
```

```
In [16]: print(X.shape, X_train.shape, X_test.shape)
```

(506, 13) (404, 13) (102, 13)

Model Training

XGBoost Regressor

```
In [17]: # Loading the model  
model = XGBRegressor()
```

```
In [18]: # training the model with X_train  
model.fit(X_train, Y_train)
```

```
Out[18]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,  
                      importance_type='gain', interaction_constraints='',  
                      learning_rate=0.300000012, max_delta_step=0, max_depth=6,  
                      min_child_weight=1, missing=nan, monotone_constraints='()',  
                      n_estimators=100, n_jobs=8, num_parallel_tree=1, random_state=0,  
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,  
                      tree_method='exact', validate_parameters=1, verbosity=None)
```

Evaluation

Prediction on training data

```
In [19]: # accuracy for prediction on training data  
training_data_prediction = model.predict(X_train)
```

```
In [20]: print(training_data_prediction)
```

```
[23.147501 20.99463 20.090284 34.69053 13.903663 13.510157
21.998634 15.1940975 10.899711 22.709627 13.832816 5.592794
29.810236 49.99096 34.89215 20.607384 23.351097 19.23555
32.695698 19.641418 26.991022 8.401829 46.00729 21.708961
27.062933 19.321356 19.288303 24.809872 22.61626 31.70493
18.542515 8.697379 17.395294 23.700663 13.304856 10.492197
12.688369 25.016556 19.67495 14.902088 24.193798 25.007143
14.900281 16.995798 15.6009035 12.699232 24.51537 14.999952
50.00104 17.525454 21.184624 31.998049 15.613355 22.89754
19.325378 18.717896 23.301125 37.222923 30.09486 33.102703
21.00072 49.999332 13.405827 5.0280113 16.492886 8.405072
28.64328 19.499939 20.586452 45.402164 39.79833 33.407326
19.83506 33.406372 25.271482 50.001534 12.521657 17.457413
18.61758 22.602625 50.002117 23.801117 23.317268 23.087355
41.700035 16.119293 31.620516 36.069206 7.0022025 20.3827
19.996452 11.986318 25.023014 49.970123 37.881588 23.123034
41.292133 17.596548 16.305374 30.034231 22.860699 19.810343
17.098848 18.898268 18.96717 22.606049 23.141363 33.183487
15.010934 11.693824 18.78828 20.80524 17.99983 19.68991
50.00332 17.207317 16.404053 17.520426 14.593481 33.110855
14.508482 43.821655 34.939106 20.381636 14.655634 8.094332
11.7662115 11.846876 18.69599 6.314154 23.983706 13.084503
19.603905 49.989143 22.300608 18.930315 31.197134 20.69645
32.21111 36.15102 14.240763 15.698188 49.99381 20.423601
16.184978 13.409128 50.01321 31.602146 12.271495 19.219482
29.794909 31.536846 22.798779 10.189648 24.08648 23.710463
21.991894 13.802495 28.420696 33.181534 13.105958 18.988266
26.576572 36.967175 30.794083 22.77071 10.201246 22.213818
24.483162 36.178806 23.09194 20.097307 19.470194 10.786644
22.671095 19.502405 20.109184 9.611871 42.799637 48.794792
13.097208 20.28583 24.793974 14.110478 21.701134 22.217012
33.003544 21.11041 25.00658 19.122992 32.398567 13.605098
15.1145315 23.088867 27.474783 19.364998 26.487135 27.499458
28.697094 21.21718 18.703201 26.775208 14.010719 21.692347
18.372562 43.11582 29.081839 20.289959 23.680176 18.308306
17.204844 18.320065 24.393475 26.396057 19.094141 13.3019905
22.15311 22.185797 8.516214 18.894428 21.792608 19.331121
18.197924 7.5006843 22.406403 20.004215 14.412416 22.503702
28.53306 21.591028 13.810223 20.497831 21.898977 23.104464
49.99585 16.242056 30.294561 50.001595 17.771557 19.053703
10.399217 20.378187 16.49973 17.183376 16.70228 19.495337
30.507633 28.98067 19.528809 23.148346 24.391027 9.521643
23.886024 49.995125 21.167099 22.597813 19.965279 13.4072275
19.948694 17.087479 12.738807 23.00453 15.222122 20.604322
26.207253 18.09243 24.090246 14.105 21.689667 20.08065
25.010437 27.874954 22.92366 18.509727 22.190847 24.004797
14.788686 19.89675 24.39812 17.796036 24.556297 31.970308
17.774675 23.356768 16.134794 13.009915 10.98219 24.28906
15.56895 35.209793 19.605724 42.301712 8.797891 24.400295
14.086652 15.408639 17.301126 22.127419 23.09363 44.79579
17.776684 31.50014 22.835577 16.888603 23.925127 12.097476
38.685944 21.388391 15.98878 23.912495 11.909485 24.960499
7.2018585 24.696215 18.201897 22.489008 23.03332 24.260433
17.101519 17.805563 13.493165 27.105328 13.311978 21.913465
20.00738 15.405392 16.595737 22.301016 24.708412 21.422579
22.878702 29.606575 21.877811 19.900253 29.605219 23.407152
13.781474 24.454706 11.897682 7.2203646 20.521074 9.725295]
```

48.30087	25.19501	11.688618	17.404732	14.480284	28.618876
19.397131	22.468653	7.0117908	20.602013	22.970919	19.719397
23.693787	25.048244	27.977154	13.393578	14.513882	20.309145
19.306028	24.095829	14.894031	26.382381	33.298378	23.61644
24.591206	18.514652	20.900269	10.406055	23.303423	13.092017
24.675085	22.582184	20.502762	16.820635	10.220605	33.81239
18.608067	49.999187	23.775583	23.909609	21.192276	18.805798
8.502987	21.50807	23.204473	21.012218	16.611097	28.100965
21.193024	28.419638	14.294126	49.99958	30.988504	24.991066
21.433628	18.975573	28.991457	15.206939	22.817244	21.765755
19.915497	23.7961				

```
In [21]: # R squared error
score_1 = metrics.r2_score(Y_train, training_data_prediction)


# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_train, training_data_prediction)

print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)
```

R squared error : 0.9999948236320982
Mean Absolute Error : 0.0145848437110976

Visualizing the actual Prices and predicted prices

```
In [22]: plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Preicted Price")
plt.show()
```



Prediction on Test Data

```
In [23]: # accuracy for prediction on test data
test_data_prediction = model.predict(X_test)
```

```
In [24]: # R squared error
score_1 = metrics.r2_score(Y_test, test_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_test, test_data_prediction)

print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)
```

```
R squared error :  0.8711660369151691
Mean Absolute Error :  2.2834744154238233
```

```
In [ ]:
```